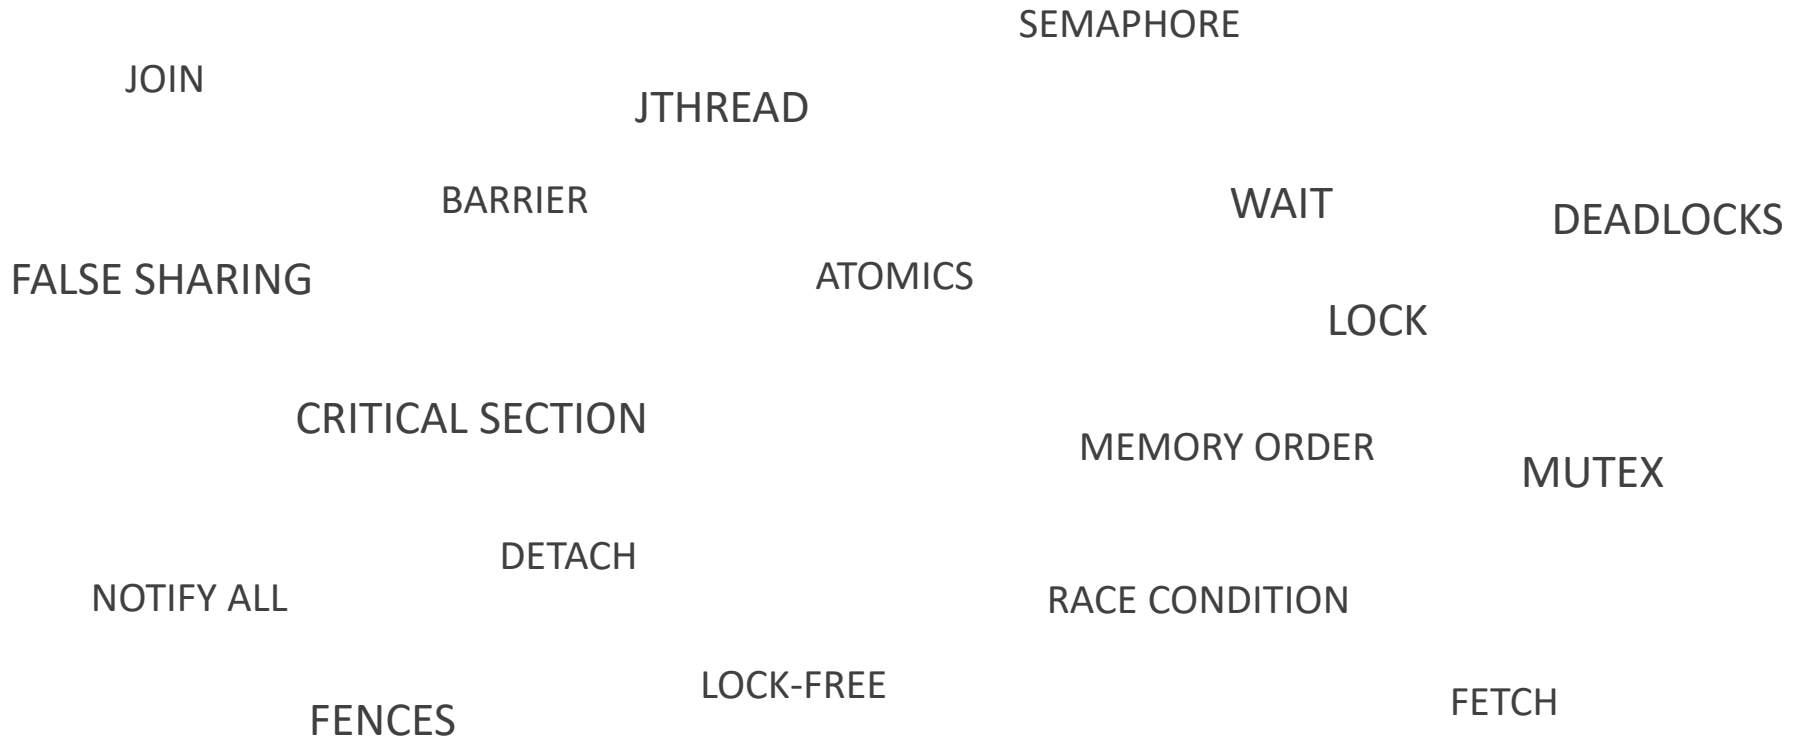


Stories from a parallel universe

Jana Machutová

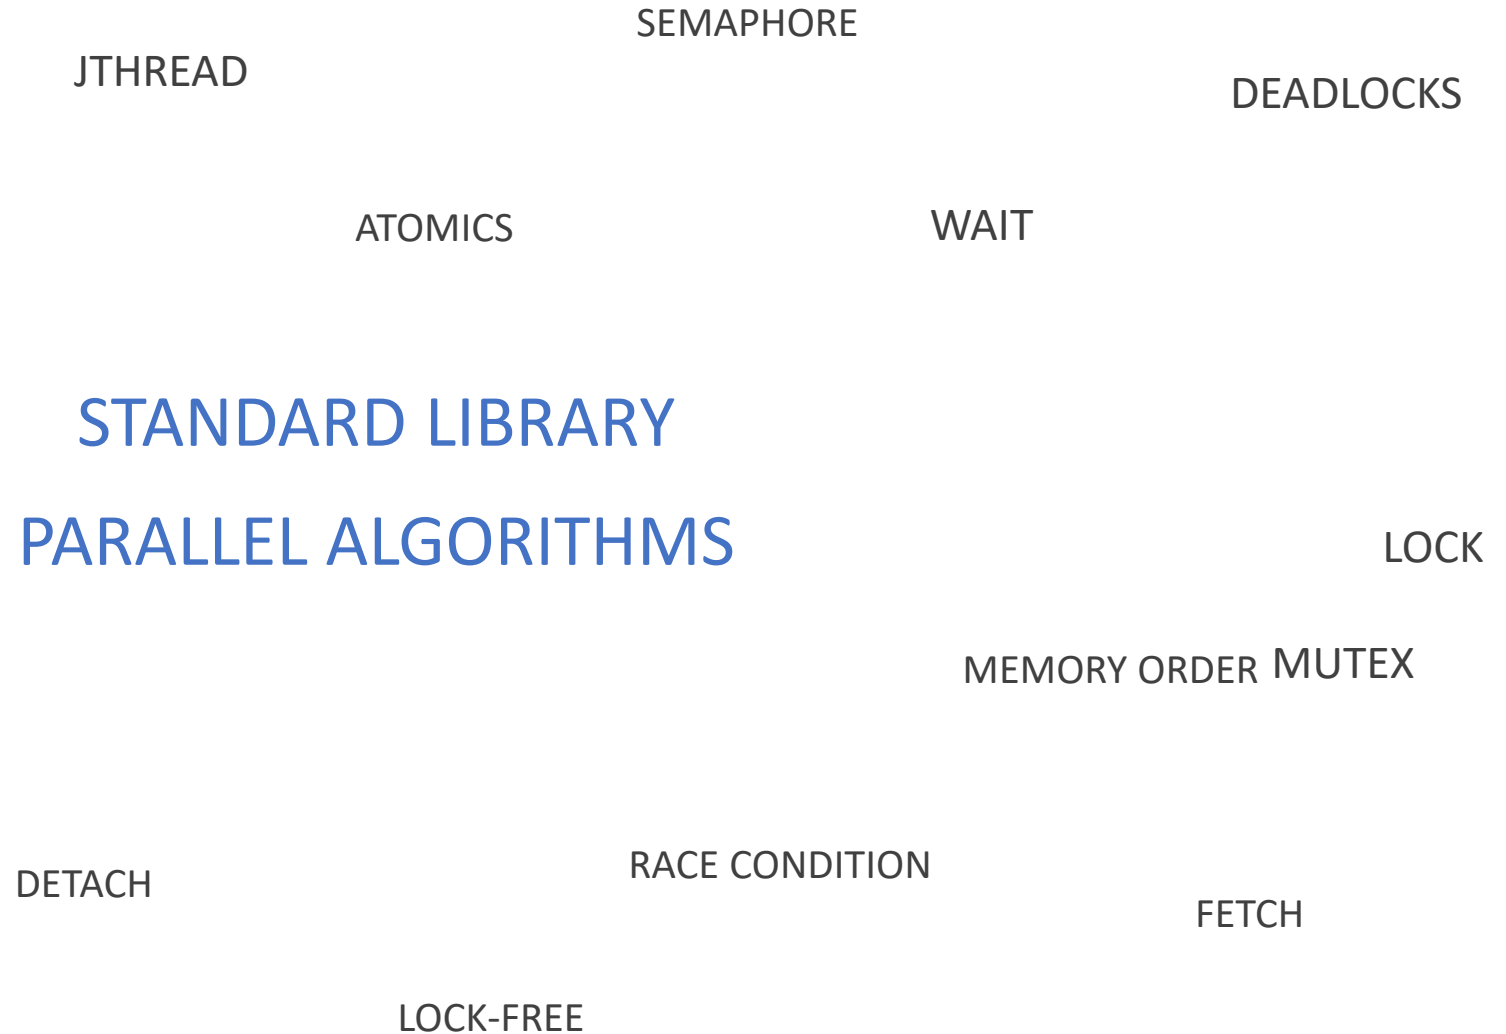
Do you feel sudden tension in your brain?



A word cloud of computer science and concurrency terms. The words are arranged in a roughly circular pattern, with some terms appearing more frequently or in larger fonts than others. The terms include:

- JOIN
- JTHREAD
- SEMAPHORE
- BARRIER
- WAIT
- DEADLOCKS
- FALSE SHARING
- ATOMICS
- LOCK
- CRITICAL SECTION
- MEMORY ORDER
- MUTEX
- DETACH
- RACE CONDITION
- NOTIFY ALL
- LOCK-FREE
- FENCES
- FETCH

We have a painkiller for you!



Standard library parallel algorithm

- Provides parallel version of most of standard library algorithms
- Introduced execution policy parameter – all the rest is same as before
- Easy switch between sequential and few types of parallel executions
- Well optimized solution to common problems
- No need to reinvent parallel wheels
- In standard since version 17 (with some changes in 20 and another on their way for future versions)
- Several new algorithms appeared
- Several old don't have any parallel version
- Not available for ranges so far

Create new worlds

No need to take care of any threads at all – it's handled automatically
Underlying thread management saves price of creating/destroying threads,
split the load equally and don't create crazy amount of threads

```
array<animal_species, 1000000> species{...};

vector<jthread> sky_residents;
for_each(species.begin(), species.end(),
    [&](const auto name){
        sky_residents.push_back(jthread([name]{
            fly_high(name);
        }));
    });
```

```
for_each(execution::par_unseq, species.begin(), species.end(),
    [](const auto name){
        swim_deep(name);
    });
```

Create new worlds

```
array<animal_species, 1000000> species{...};
```

```
vector<jthread> sky_residents;  
for_each(species.begin(), species.end(),  
    [&](const auto name){  
        sky_residents.push_back(jthread([name]{  
            fly_high(name);  
        }));  
    });
```

```
for_each(execution::par_unseq, species.begin(), species.end(),  
    [](const auto name){  
        swim_deep(name);  
    });
```

Thread #0 species [0]

Thread #1 species [1]

**RESOURCE
UNAVAILABLE**

Thread #999 999 species [999 999]

Thread #0 species [0 – 9 999]

Thread #1 species [10 000 – 19 999]

Thread #100 species [90 000 – 99 999]

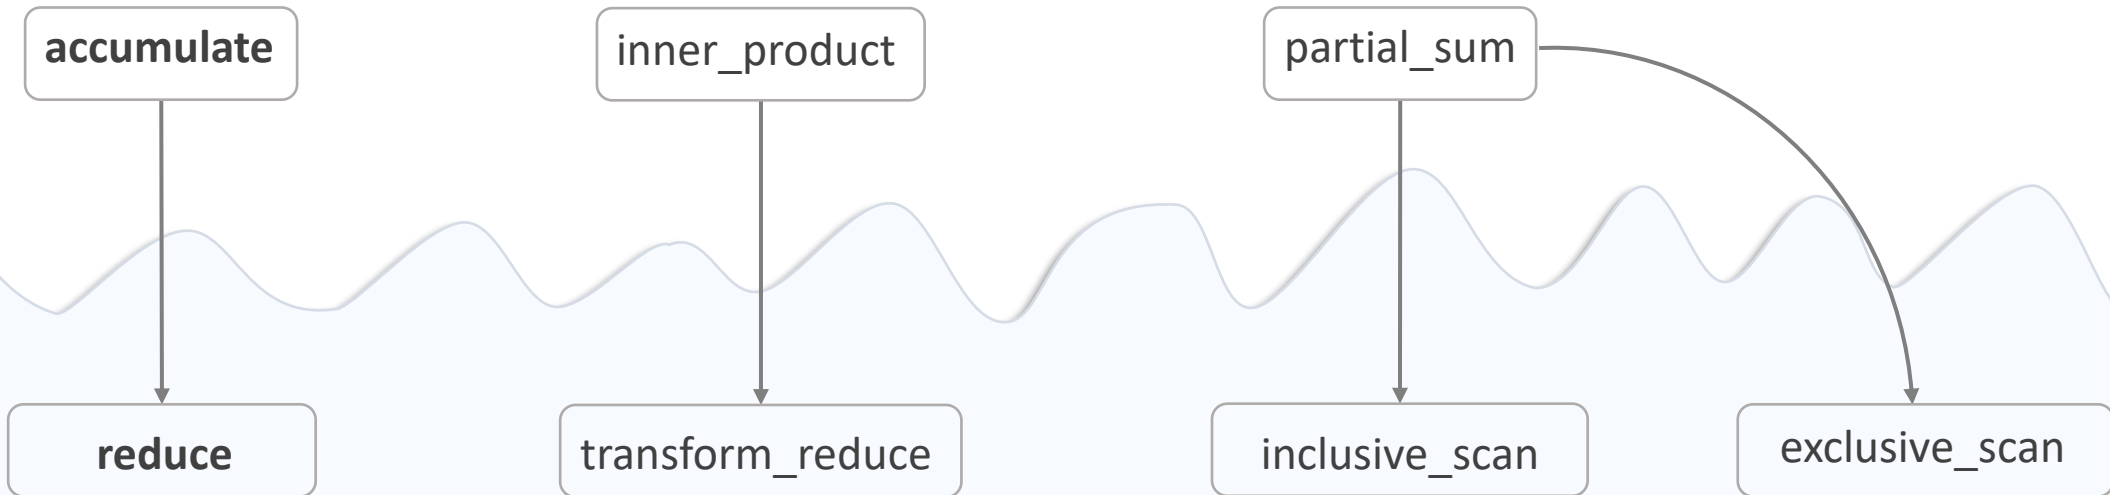
Why do you want to use parallel algorithms

- No need for direct manipulation with threads
- No need to care about threads lifetime
- Clever thread management is already part of the package
- Widely tested and optimized solution
- Easy to switch between different execution policies
- Improved readability
- Less error prone
- If you can use sequential version, you will be able to use parallel without any additional studies

Do not reinvent a wheel – use standard library as much as you can

New algorithms introduced

All of them are here as an alternative to existing sequential only solution



accumulate vs reduce

First sight the only difference is possibility to add policy

Second sight try to modify operation to some custom

```
vector<int> species_count{68, 15, 4, 45, 18, 3, 2, 11};
```

```
accumulate(species_count.begin(), species_count.end(), 0, plus{});
```

output: **166**

```
reduce(execution::par, species_count.begin(), species_count.end(), 0 , plus{});
```

output: **166**

accumulate vs reduce

How many animals will be there in next generation with growth_factor 2

Non-associative reduction function makes mess

Have a look what's going on in both cases on next slides

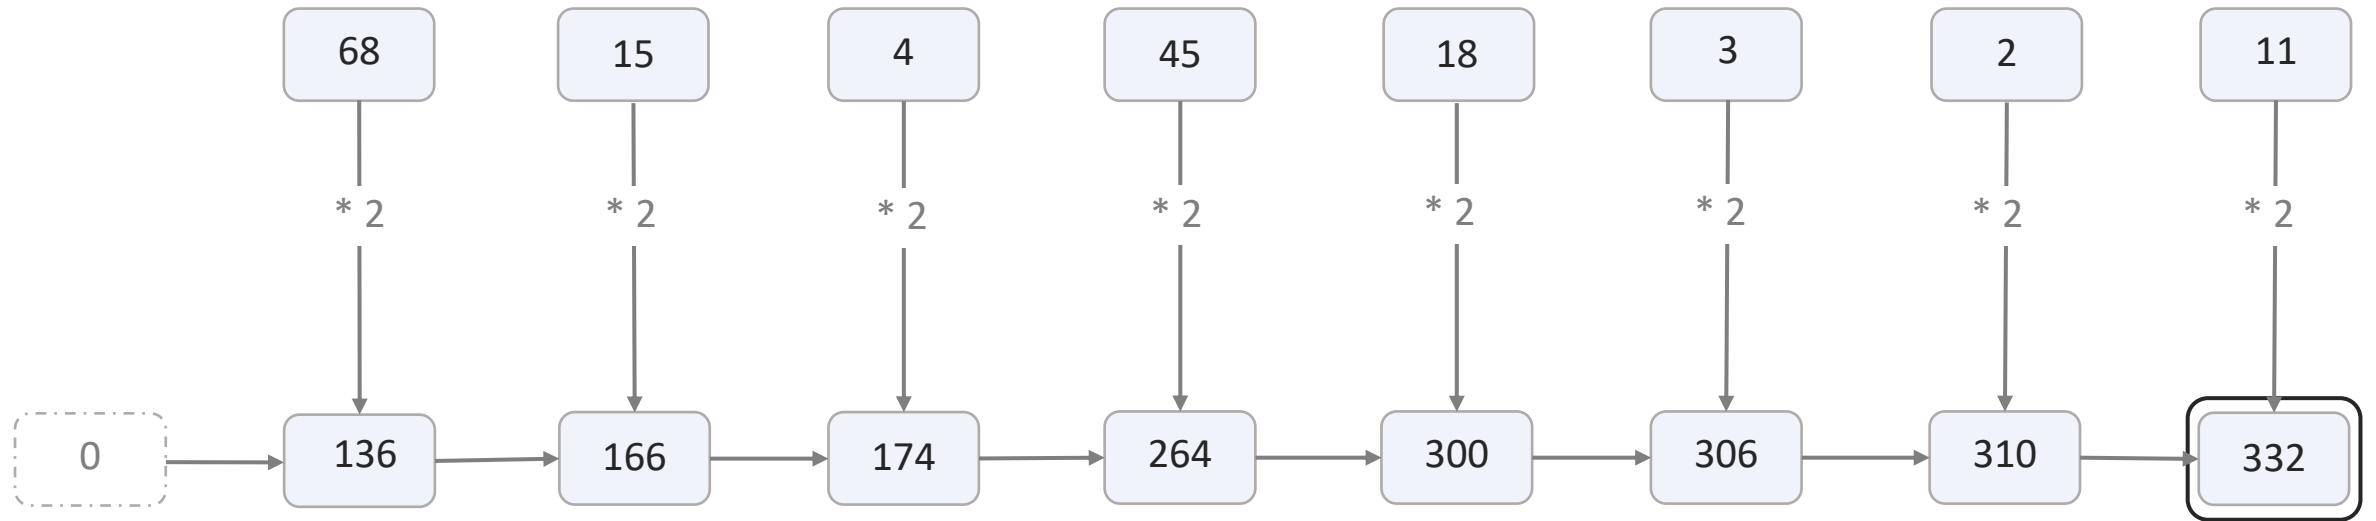
```
vector<int> species_count{68, 15, 4, 45, 18, 3, 2, 11};  
int growth_factor = 2;  
  
accumulate(species_count.begin(), species_count.end(), 0,  
           [growth_factor](const auto first, const auto second){  
               return first + second * growth_factor;  
           }));
```

output: **332**

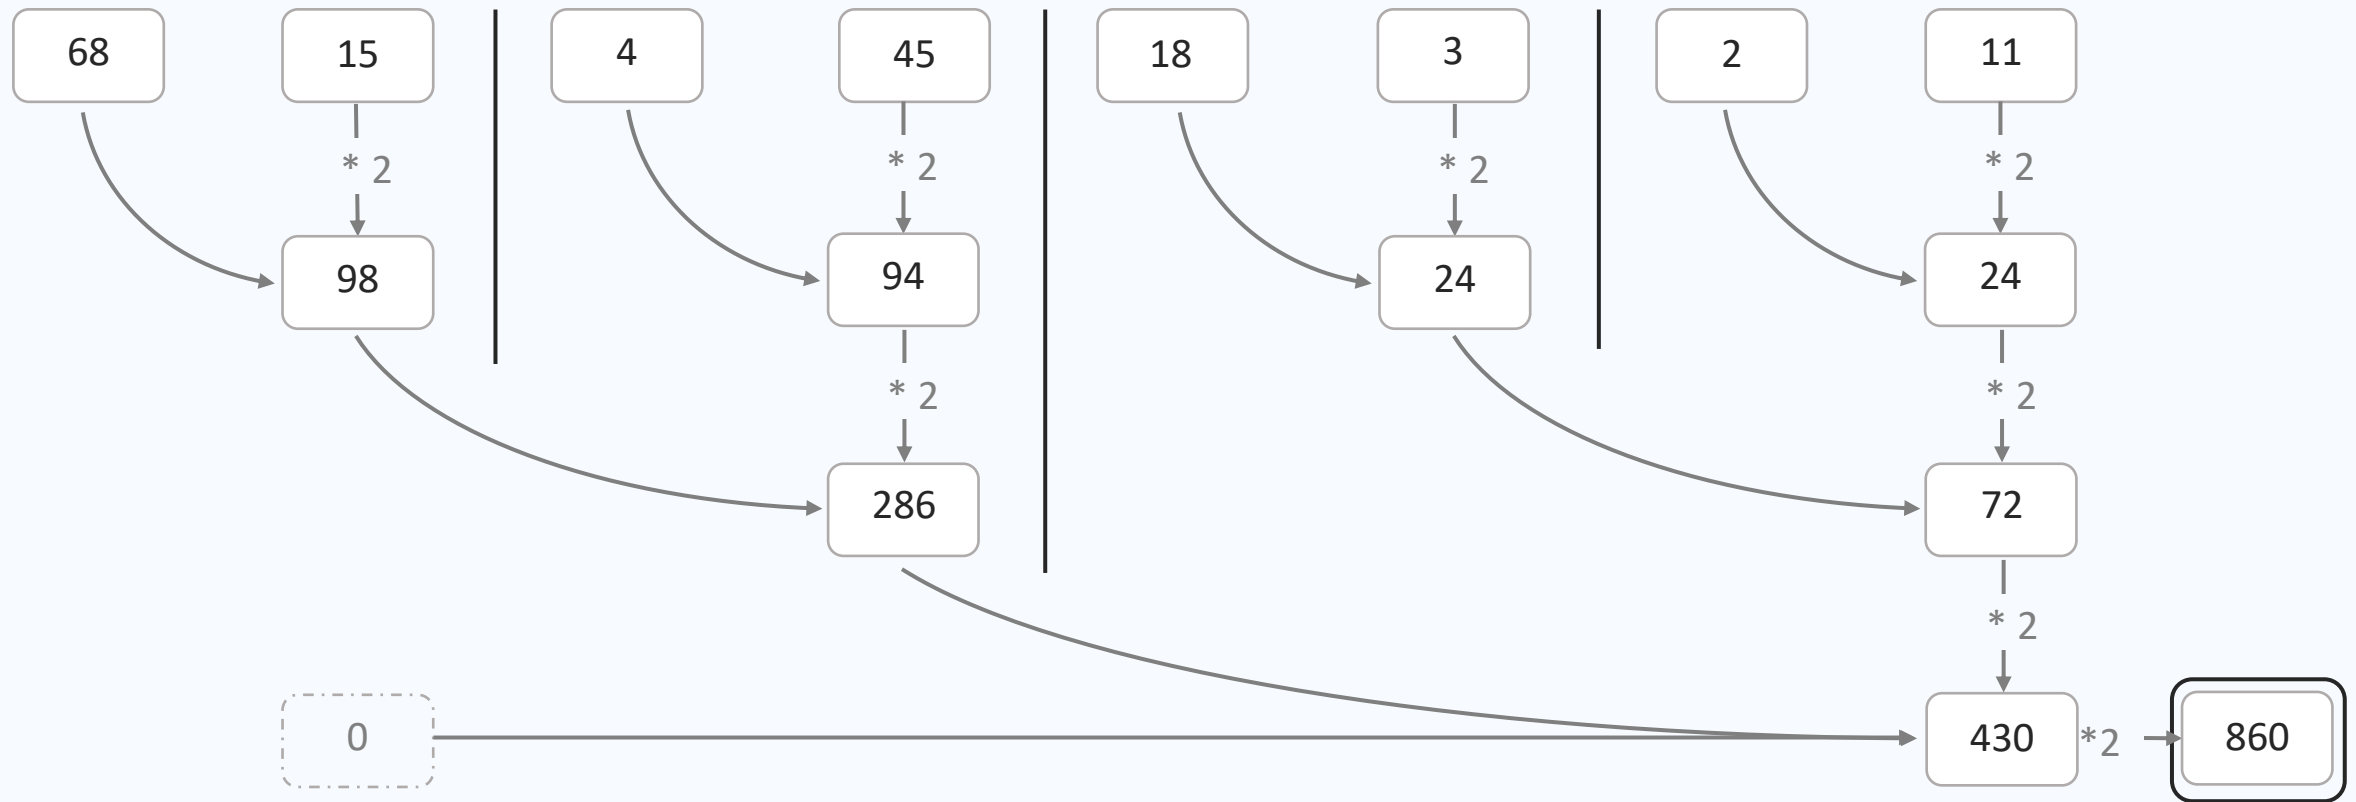
```
reduce(execution::par, species_count.begin(), species_count.end(), 0,  
       [growth_factor](const auto first, const auto second){  
           return first + second * growth_factor;  
       }));
```

output: **716**

accumulate walktrough



reduce walktrough



$$0 + (68 + 2*15 + 2*(4 + 2*45) + 2*(18 + 2*3 + 2*(2 + 2*11))) * 2 \quad != \quad 2*68 + 2*15 + 2*4 + 2*45 + 2*18 + 2*3 + 2*2 + 2*11$$

accumulate vs reduce

```
vector<int> species_count{68, 15, 4, 45, 18, 3, 2, 11};  
int growth_factor = 2;  
  
accumulate(species_count.begin(), species_count.end(), 0,  
           [growth_factor](const auto first, const auto second){  
               return first + second * growth_factor;  
           });
```

output: **322**

```
reduce(execution::par, species_count.begin(), species_count.end(), 0,  
       [growth_factor](const auto first, const auto second){  
           return first + second * growth_factor;  
       });
```

output: **716**

accumulate vs reduce - fixed

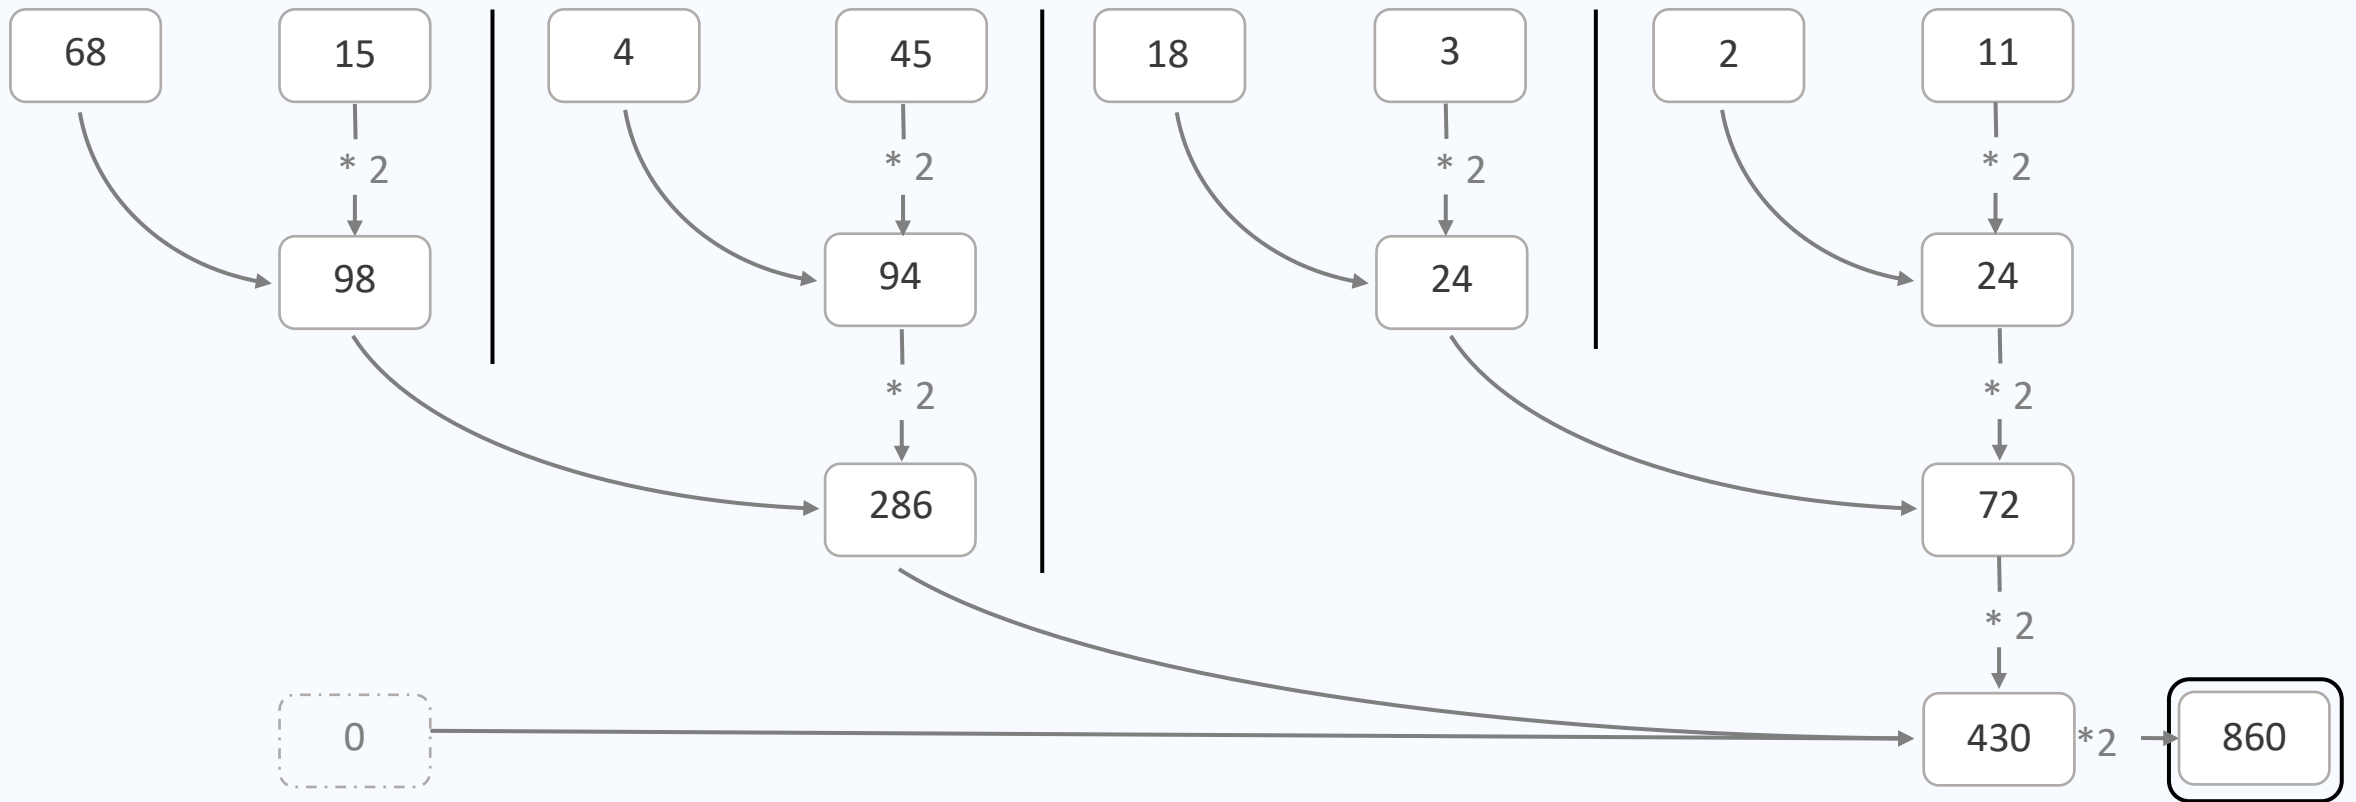
```
vector<int> species_count{68, 15, 4, 45, 18, 3, 2, 11};  
int growth_factor = 2;  
  
accumulate(species_count.begin(), species_count.end(), 0,  
           [growth_factor](const auto first, const auto second){  
               return first + second * growth_factor;  
           });
```

output: **332**

```
transform_reduce(execution::par, species_count.begin(), species_count.end(), 0, std::plus{},  
                 [growth_factor](const auto item){  
                     return item * growth_factor;  
                 });
```

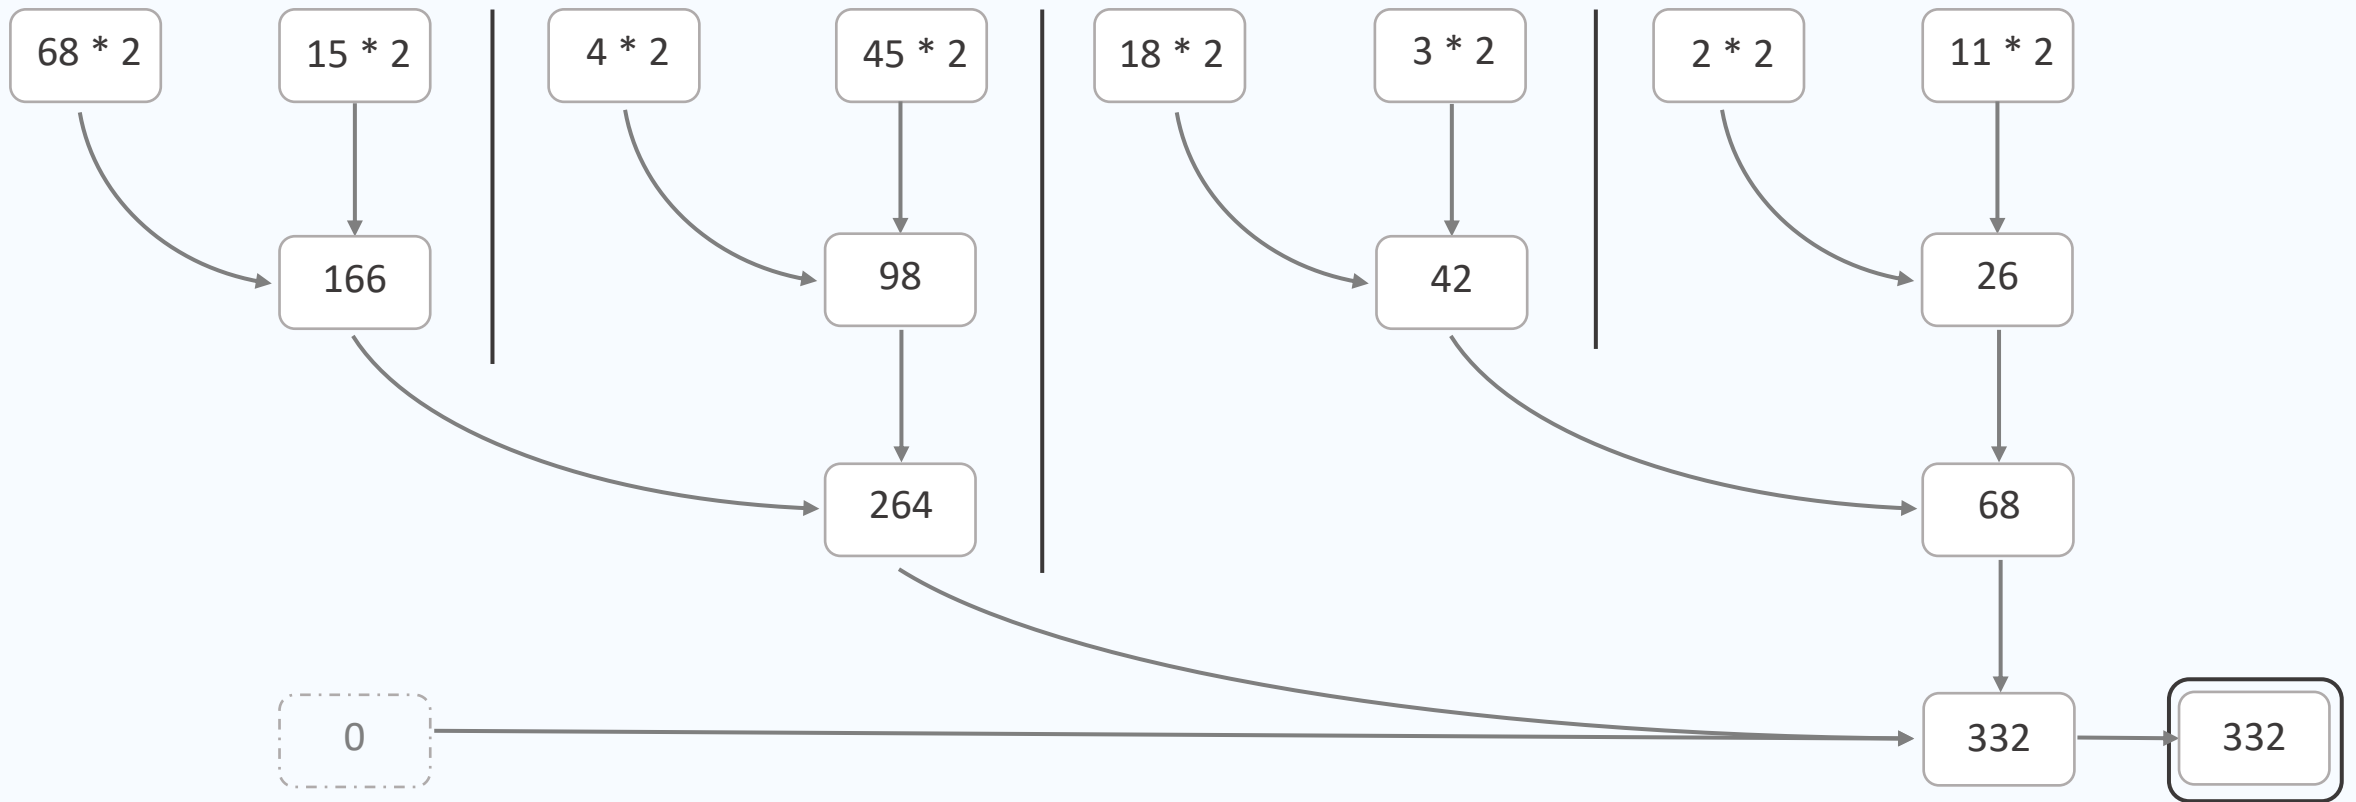
output: **332**

reduce walktrough



$$0 + (68 + 2*15 + 2*(4 + 2*45) + 2*(18 + 2*3 + 2*(2 + 2*11))) * 2 \quad != \quad 2*68 + 2*15 + 2*4 + 2*45 + 2*18 + 2*3 + 2*2 + 2*11$$

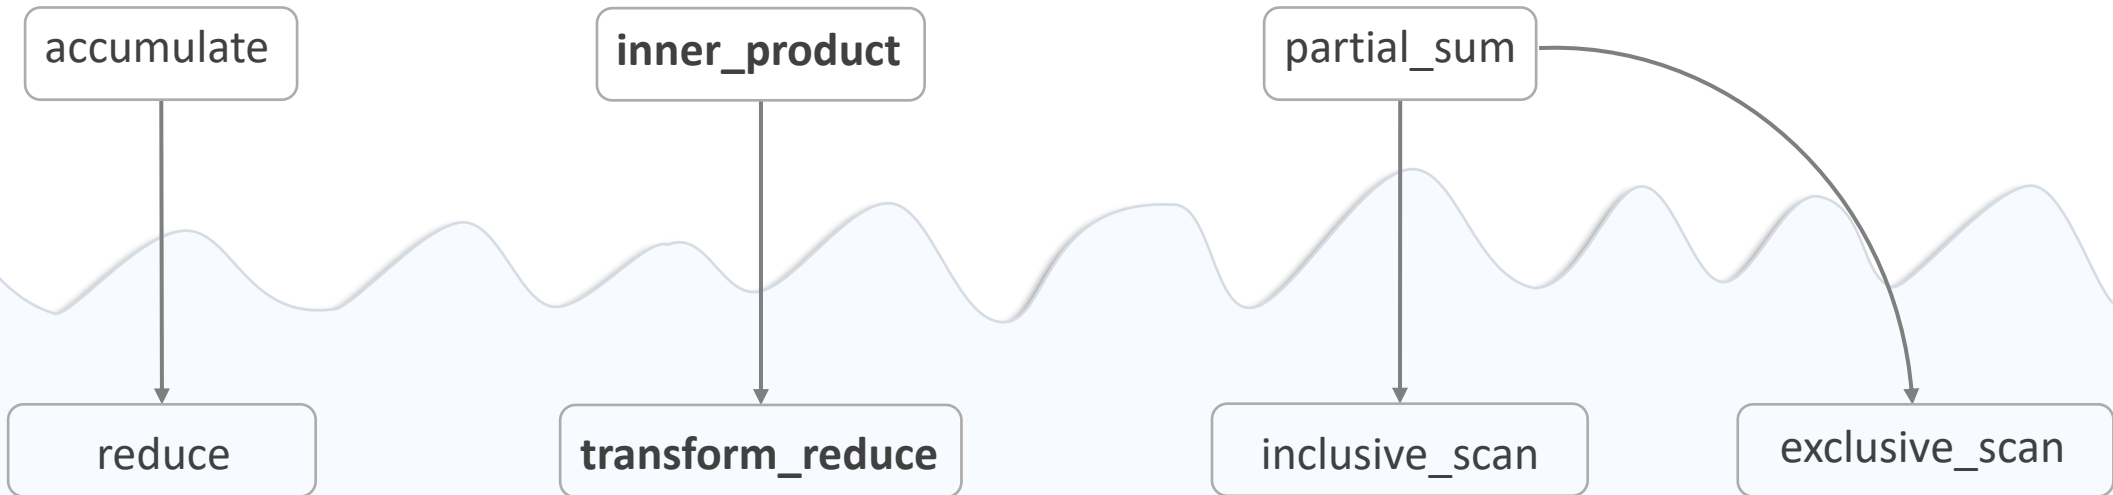
transform_reduce walkthrough



$$0 + (136 + 30 + (8 + 90) + (36 + 6 + (4 + 22))) == 0 + 2*68 + 2*15 + 2*4 + 2*45 + 2*18 + 2*3 + 2*2 + 2*11$$

New algorithms introduced

All of them are here as an alternative to existing sequential only solution



inner_product vs transform_reduce

```
vector<int> weights{1e-6f, 2.1e-5f, 3.f, 43.3f, 5.1e3f, 6.5e4f, 7.7e6f, 1.5e8f};  
vector<int> species_count{68, 15, 4, 45, 18, 3, 2, 11};
```

```
inner_product(species_count.begin(), species_count.end(), weights.begin(), 0);
```

$0 + 1e-6f*68 + 2.1e-5f*15 + 3.f*4 + 43.3f*45 + 5.1e3f*18 + 6.5e4f*3 + 7.7e6f*2 + 1.5e8f*11$

output: **1 665 688 704 g**

output with double: **1 665 688 761 g**

```
transform_reduce(execution::par, species_count.begin(), species_count.end(), weights.begin(), 0);
```

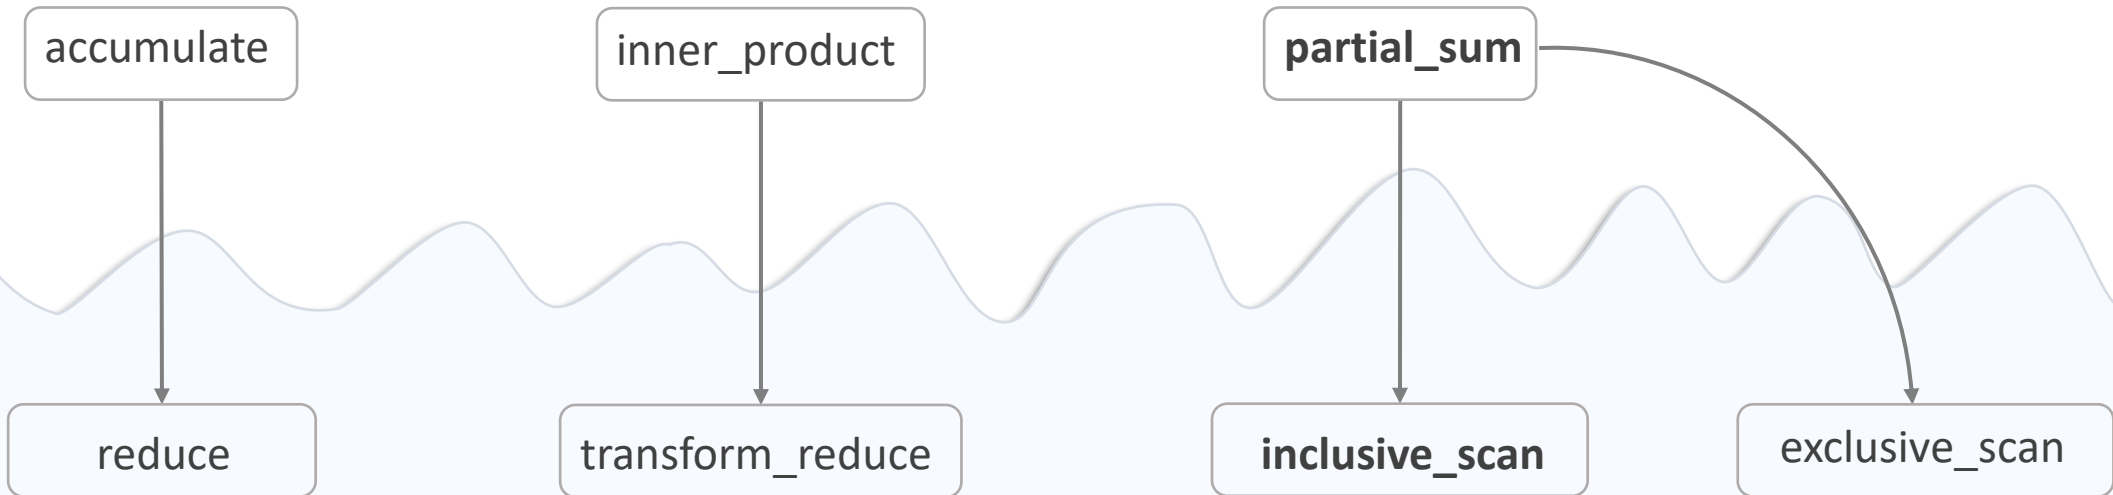
$0 + (1e-6f*68 + 2.1e-5f*15 + (3.f*4 + 43.3f*45) + (5.1e3f*18 + 6.5e4f*3 + (7.7e6f*2 + 1.5e8f*11)))$

output: **1 665 688 832 g**

output with double: **1 665 688 761 g**

New algorithms introduced

All of them are here as an alternative to existing sequential only solution



partial_sum vs inclusive_scan

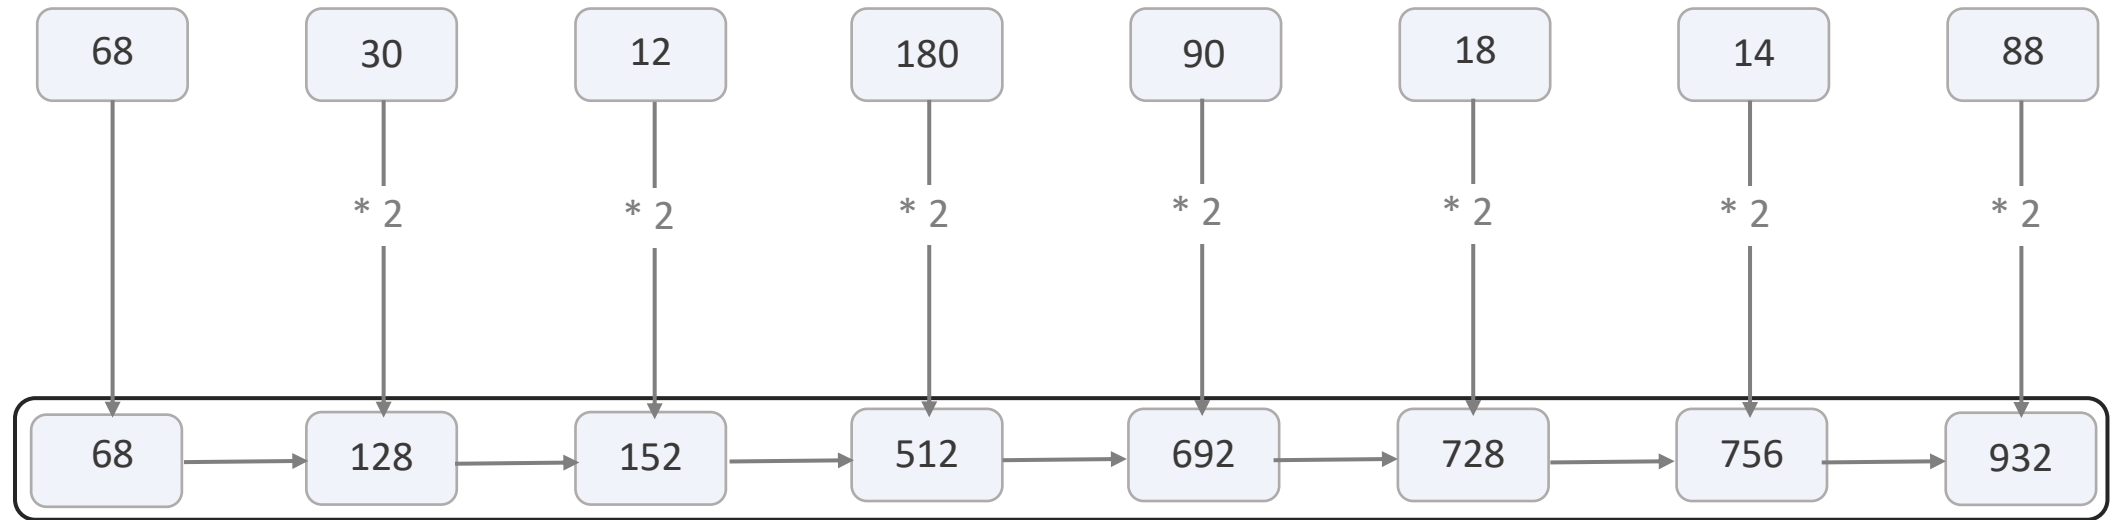
```
vector<int> species_mass{68, 30, 12, 180, 90, 18, 14, 88};  
array<int, 8> edible_mass;  
const int growth_factor = 2;  
  
partial_sum(species_mass.begin(), species_mass.end(), edible_mass.begin(),  
            [growth_factor](const auto& smaller_pray, const auto& equal_pray){  
                return equal_pray * growth_factor + smaller_pray;  
            });
```

output: **68, 128, 152, 512, 692, 728, 756, 932**

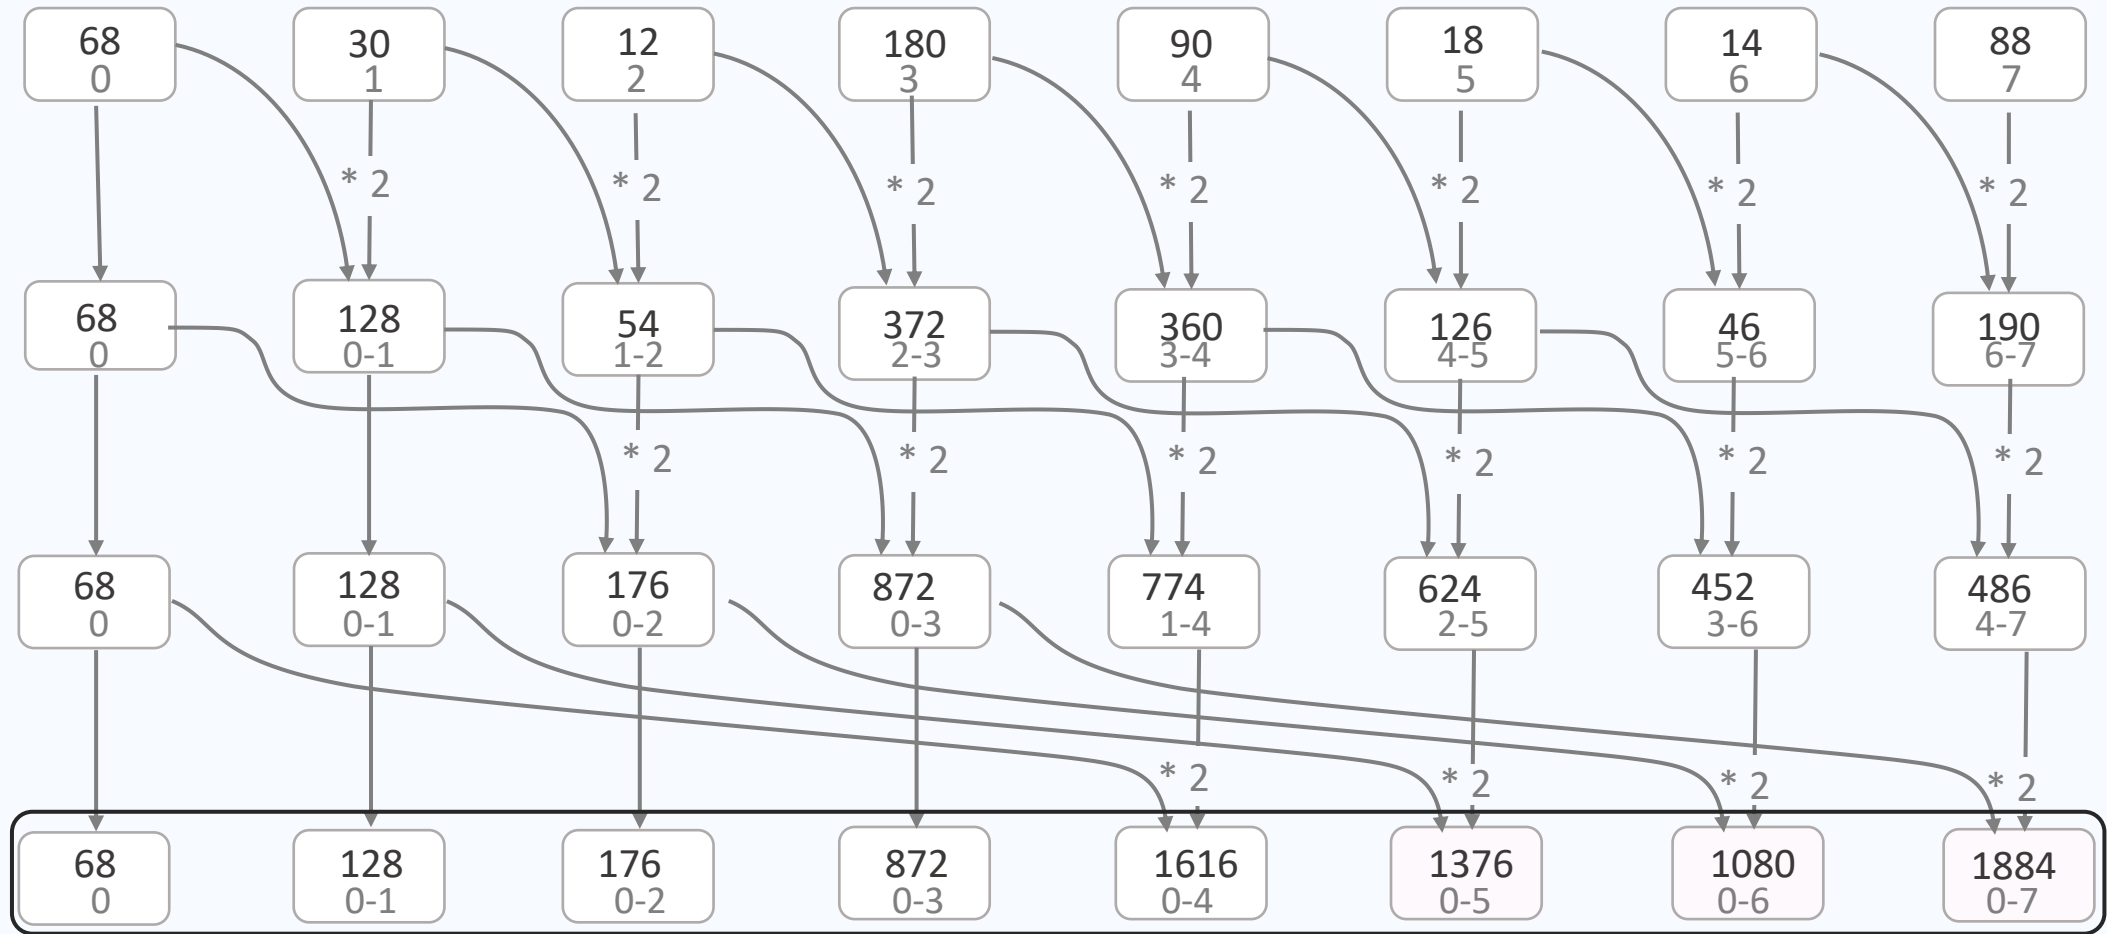
```
inclusive_scan(execution::par, species_mass.begin(), species_mass.end(), edible_mass.begin(),  
              [growth_factor](const auto& smaller_pray, const auto& equal_pray){  
                  return equal_pray * growth_factor + smaller_pray;  
              });
```

output: **68, 128, 176, 872, 1616, 1376, 1080, 1884**

partial_sum walktrough

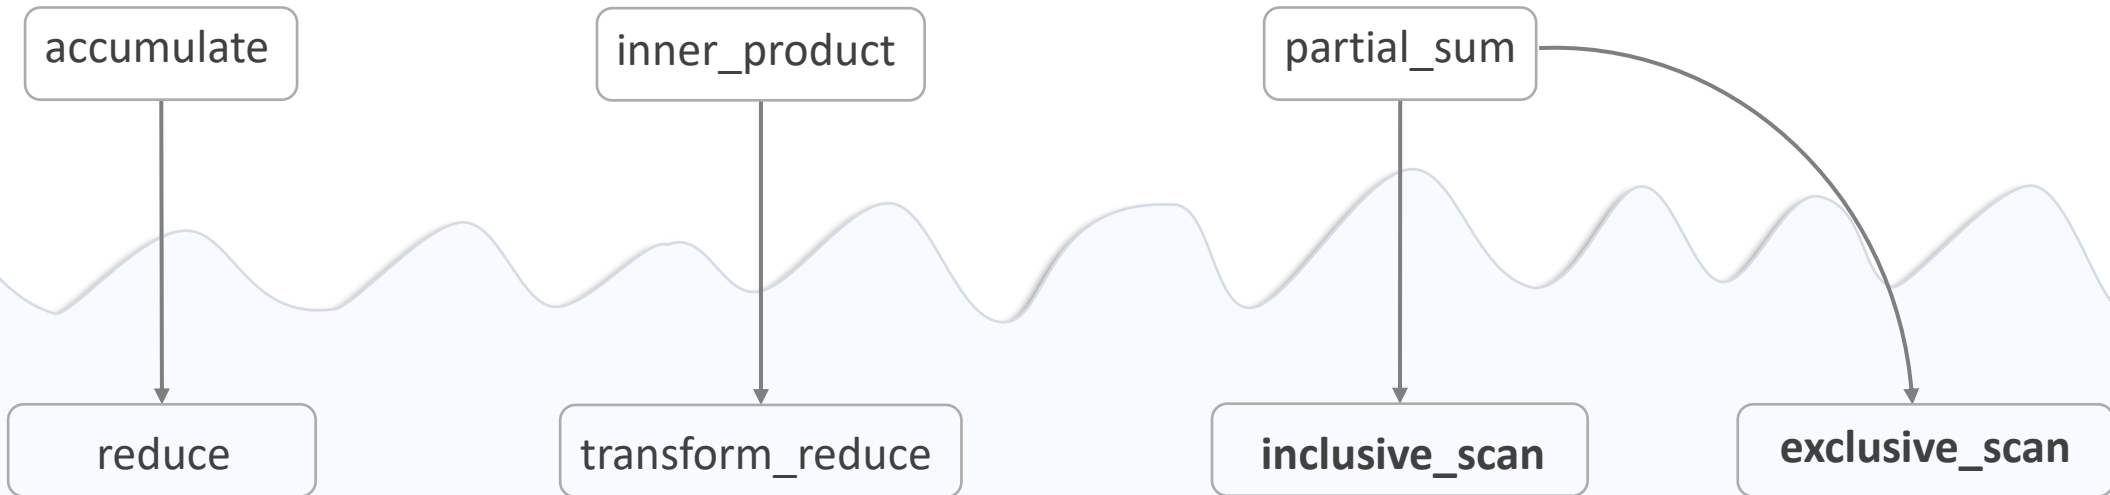


naïve inclusive_scan walkthrough



New algorithms introduced

All of them are here as an alternative to existing sequential only solution



inclusive_scan vs exclusive_scan

```
vector<int> species_mass{136, 60, 24, 360, 180, 36, 28, 176};  
array<int, 8> edible_mass;
```

```
inclusive_scan(species_mass.begin(), species_mass.end(), edible_mass.begin());
```

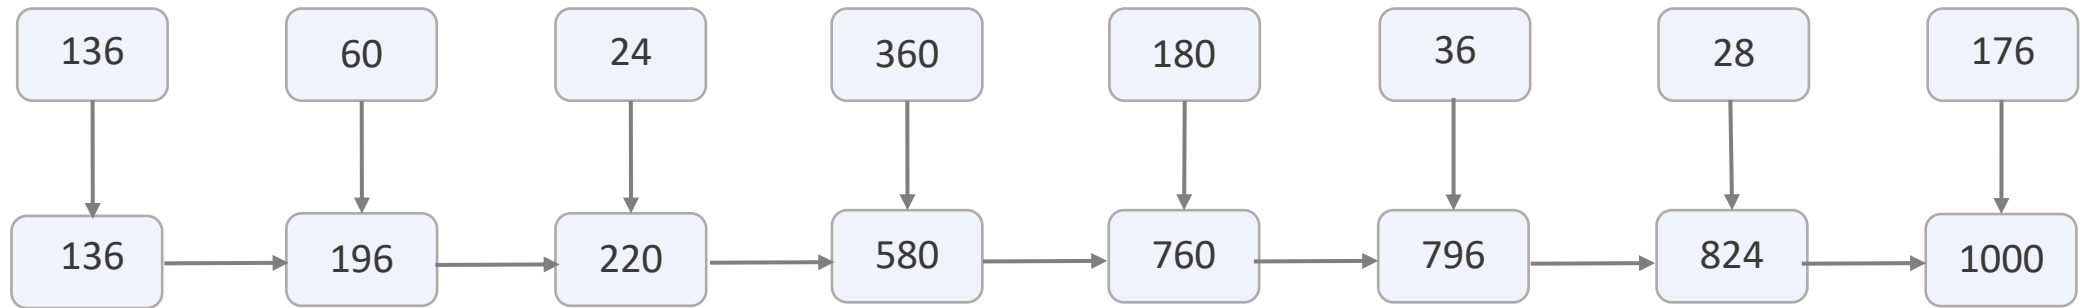
output: **136, 196, 220, 580, 760, 796, 824, 1000**

```
exclusive_scan(species_mass.begin(), species_mass.end(), edible_mass.begin(), 0);
```


output: **0, 136, 196, 220, 580, 760, 796, 824**

inclusive_scan vs exclusive_scan

inclusive_scan:




exclusive_scan:



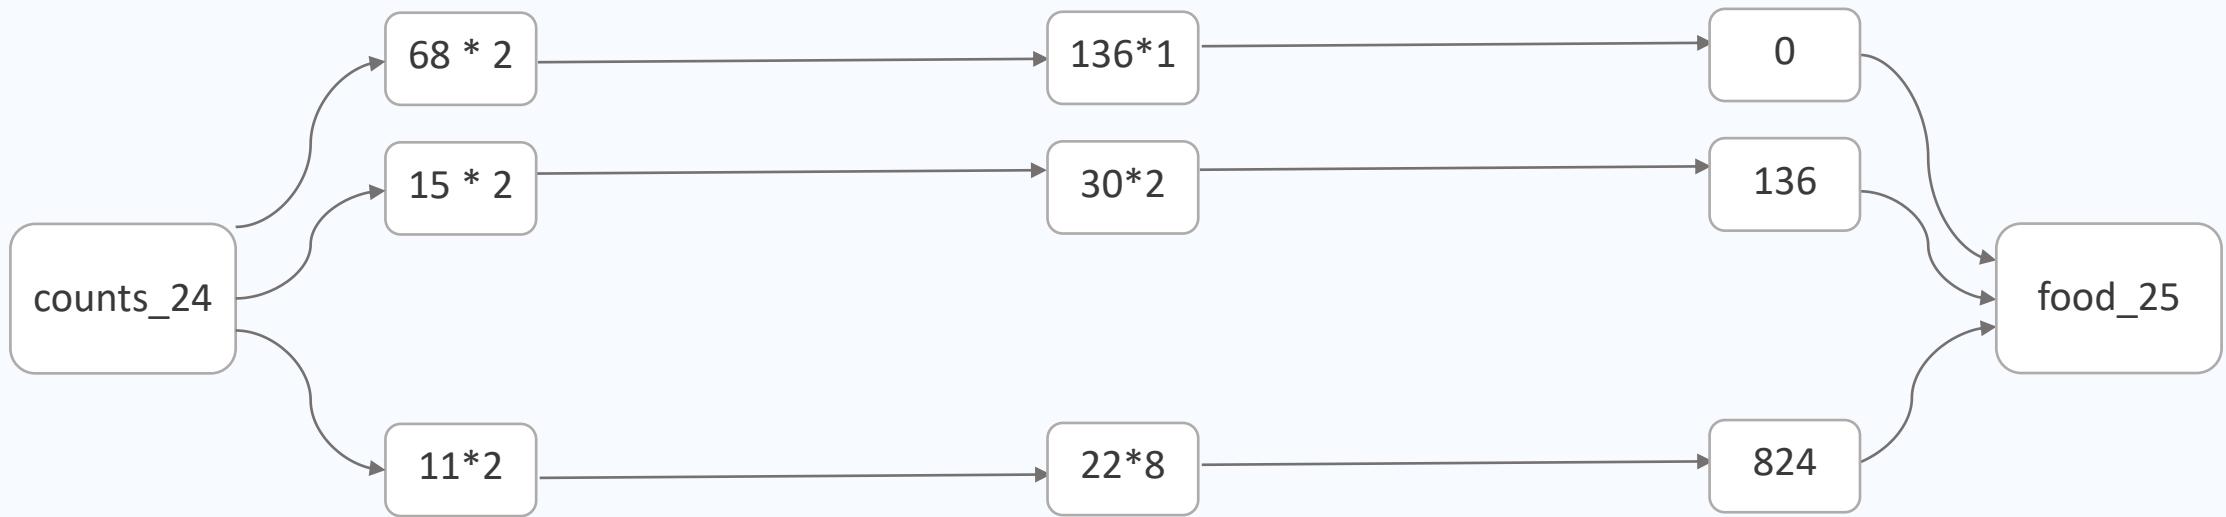
Unwanted synchronization points

```
array<int, 8> counts_24{68, 15, 4, 45, 18, 3, 2, 11};  
array<int, 8> weights{1, 2, 3, 4, 5, 6, 7, 8};
```

```
transform(par_unseq, counts_24.begin(), counts_24.end(), counts_25.begin(), [](auto i){return i * 2;});  
transform(par_unseq, counts_25.begin(), counts_25.end(), weights.begin(), weight_25.begin(), multiplies{});  
exclusive_scan(par_unseq, weight_25.begin(), weight_25.end(), food_25.begin(), 0);
```



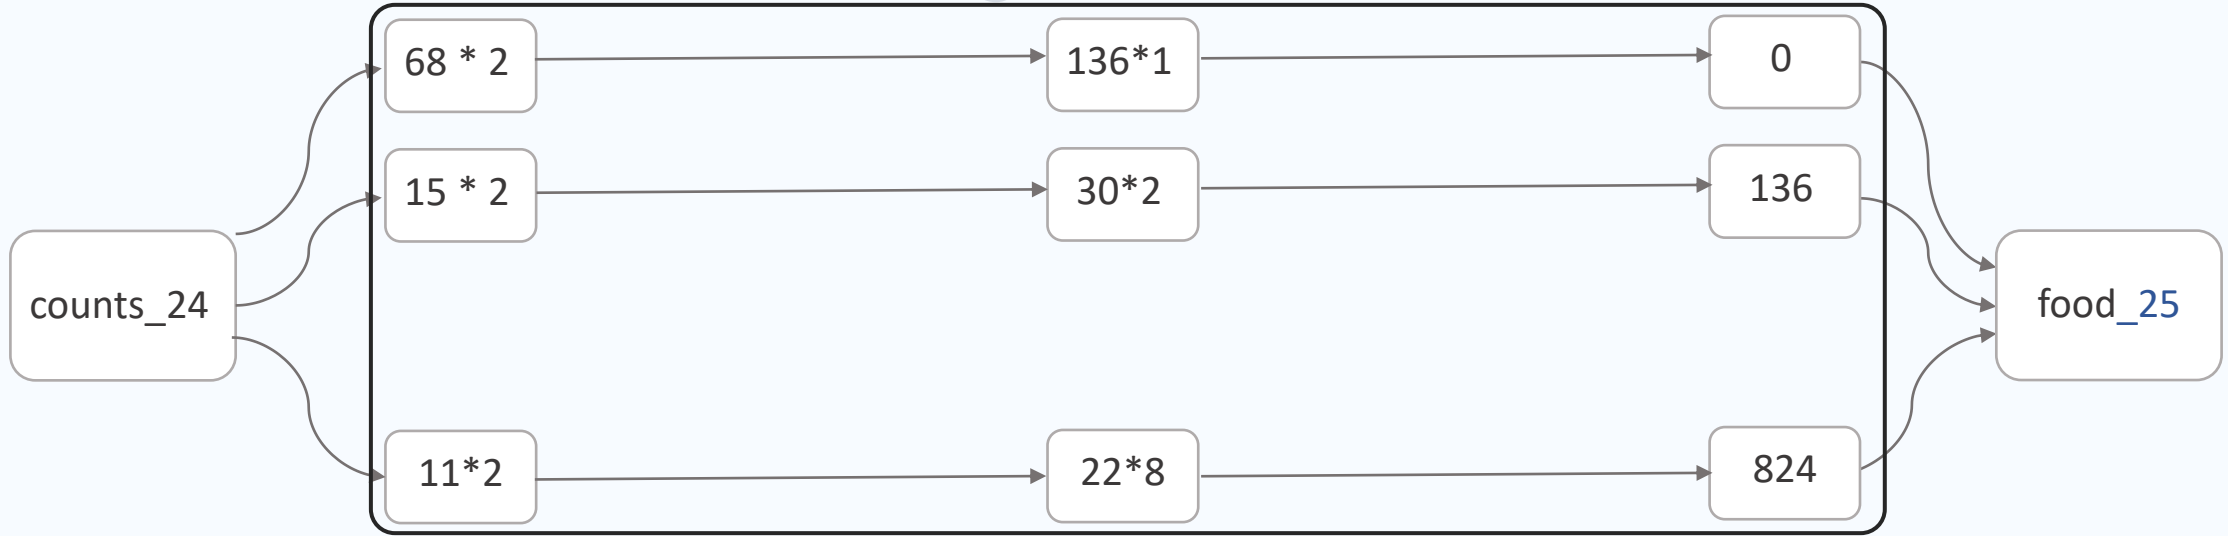
Unwanted synchronization points



```
auto weighs_25 = views::transform([&](const auto item){return item * 2;}) |  
    views::zip_transform(multiplies{}, counts_24, weighs);
```

```
exclusive_scan(par_unseq, weighs_25.begin(), weighs_25.end(), food_25.begin(), 0);
```

Unwanted synchronization points

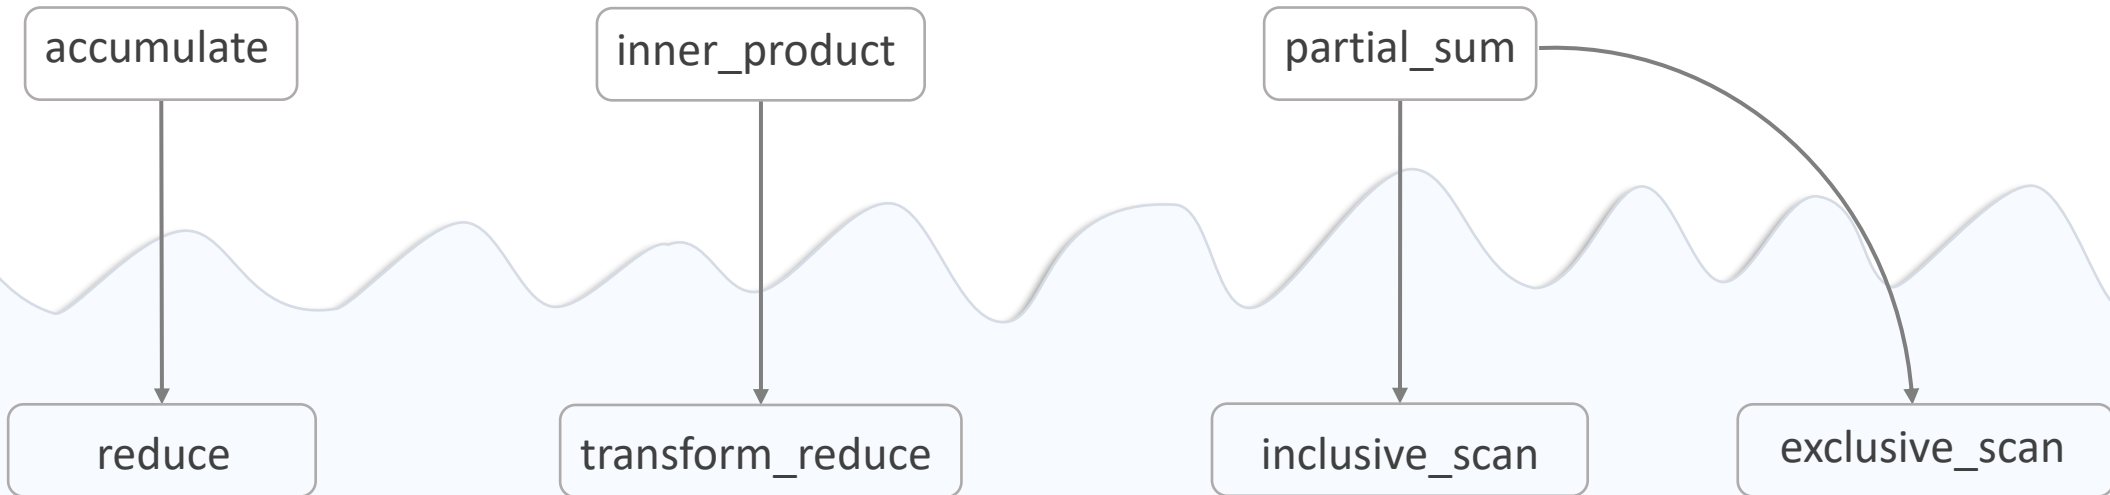


```
auto weighs_25 = views::zip_transform(multiplies{}, counts_24, weights);
```

```
transform_exclusive_scan(par_unseq, weighs_25.begin(), weighs_25.end(), food_25.begin(), 0,  
    plus{}, ([&](const auto item){return item * 2;}));
```

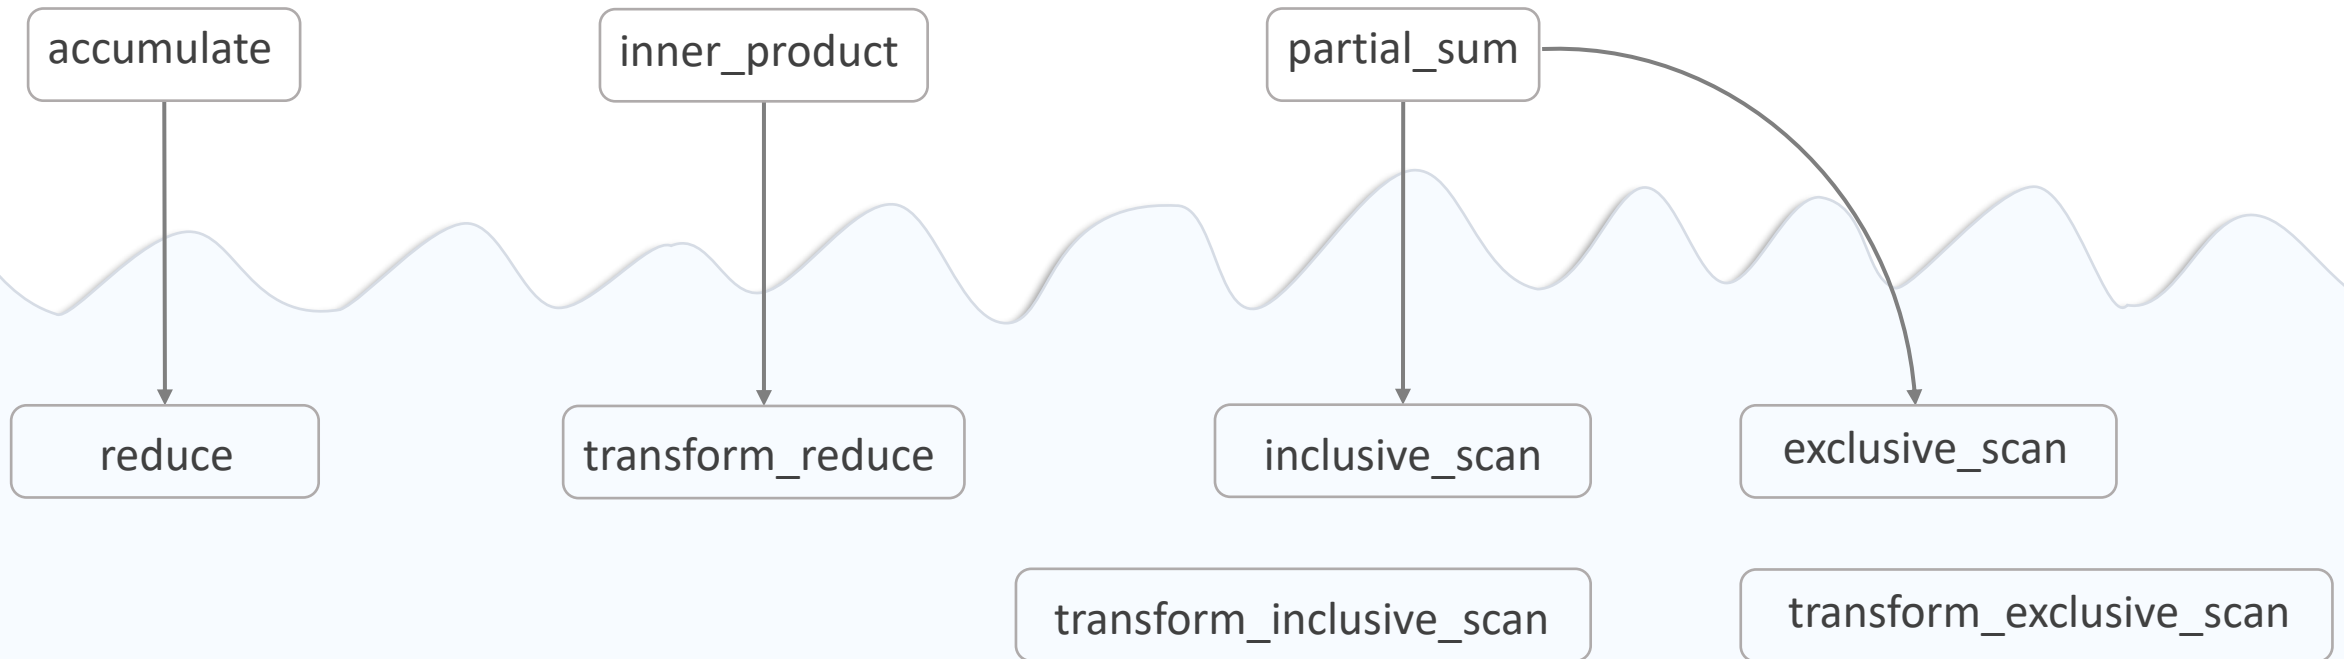
New algorithms introduced

All of them are here as an alternative to existing sequential only solution



New algorithms introduced

All of them are here as an alternative to existing sequential only solution



Execution policy

- First param of each parallel algorithm
 - Defines how much is it possible to parallelize code inside algorithm body
 - For the best results go for the strongest policy you can
 - The stronger policy you get, the more limited you are inside the algorithm body
 - If you are unable to use strongest one, maybe you won't gain from parallelization
 - Always measure performance before and after parallelization
-
- `execution::seq`
 - `execution::par`
 - `execution::unseq`
 - `execution::par_unseq`

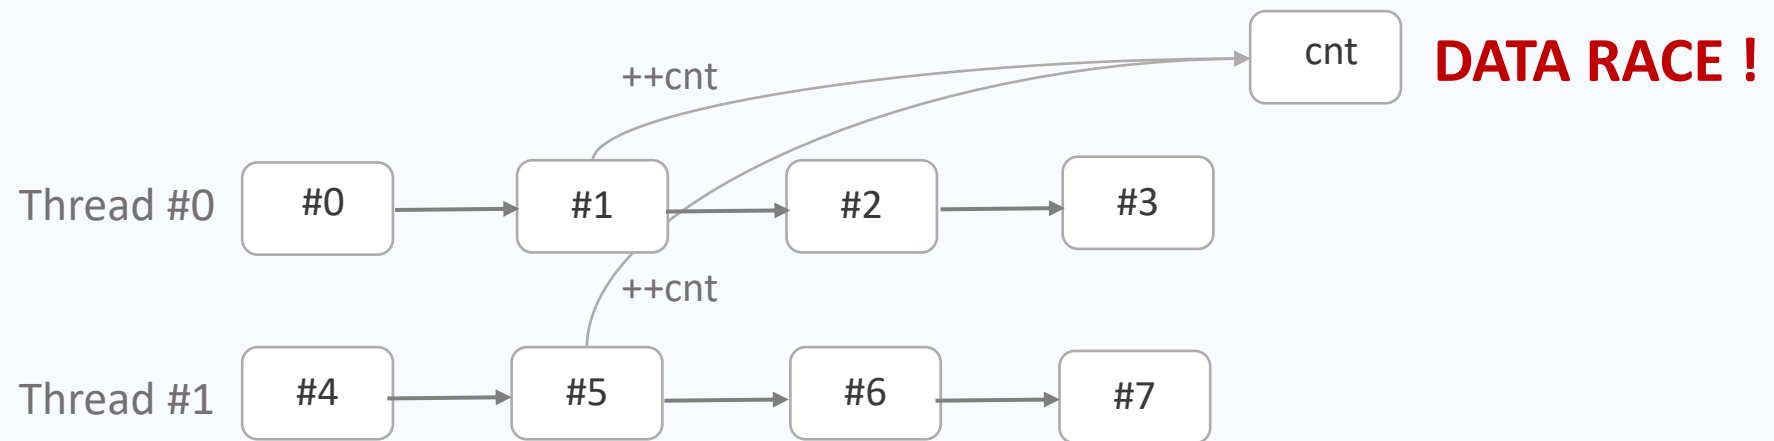
execution::seq

- Code is executed sequentially
- All computations are executed on single thread one by one
- No parallelization at all
- If there is no policy defined it runs in execution::seq




execution::par

- Code is executed on multiple threads
- Order in which are elements processed is not predictable
- Shared data needs to be protected as they may be accessed by different threads concurrently
- If blocking protection is needed it may not perform better than sequential approach – don't forget to measure it



execution::unseq

- Code is executed in single thread
- Whole vector of data is executed together – their instructions are in unpredictable order
- No synchronization is allowed
- No, not even atomics



execution::par_unseq

- Data are vectorized and executed on multiple threads
- Strongest possible version of parallelization
- Keeps the same limitation as unseq – no synchronization is allowed
- If your code may run on unseq it may run on par_unseq there is no difference in rules
- You should try to target this policy



cnt

**DON'T EVEN
THINK ABOUT IT**

Plenty of fish in the sea

Easiest interactive example

What policy could we use?

```
struct Fish{ int pos; int velocity;};  
vector<Fish> fishes{{32, 10}, {54, 15}, {55, 5}, {123, 20}, {124, 18}, {124, 18}, {320, 2}, {323, 5}, {480, 16}};  
const auto time_interval = 1;  
  
for_each(fishes.begin(), fishes.end(), [&](auto& fish)  
{  
    fish.pos += time_interval * fish.velocity;  
}));
```

Plenty of fish in the sea

```
struct Fish{ int pos; int velocity;};  
vector<Fish> fishes{{32, 10}, {54, 15}, {55, 5}, {123, 20}, {124, 18}, {124, 18}, {320, 2}, {323, 5}, {480, 16}};  
const auto time_interval = 1;  
  
for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)  
{  
    fish.pos += time_interval * fish.velocity;  
}));
```

Beware of shark

```
struct Fish{ int pos; int velocity;};
vector<Fish> fishes{{32, 10}, {54, 15}, {55, 5}, {123, 20}, {124, 18}, {124, 18}, {320, 2}, {323, 5}, {480, 16}};
const auto time_interval = 1;
bool shark = false;           //shark is handled somewhere else and this will signal his presence

for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)
{
    if(shark)
        fish.run_away();
    else
        fish.pos += time_interval * fish.velocity;
});
```

Beware of shark

Atomic shark allows us to go par, but it's synchronization so it's not allowed to go par_unseq. We can reason that this very specific scenario would be safe, but it's still not allowed by standard.

But you don't lose much. It will be impossible to vectorize atomic operations anyway, so we would end up with par anyway

```
struct Fish{ int pos; int velocity;};
vector<Fish> fishes{{32, 10}, {54, 15}, {55, 5}, {123, 20}, {124, 18}, {124, 18}, {320, 2}, {323, 5}, {480, 16}};
const auto time_interval = 1;
atomic<bool> shark = false;      //shark is handled somewhere else and this will signal his presence

for_each(execution::par, fishes.begin(), fishes.end(), [&](auto& fish)
{
    if(shark)
        fish.run_away();
    else
        fish.pos += time_interval * fish.velocity;
});
```

Look for a hideout

Again - what's wrong and how could we fix it

Fish is looking for safety.
Each hideout has just limited capacity.
Check only those with free slots.
Choose the closest on left and right
Go for the closest
Swimming takes some time
Try to get inside and hope nobody was faster

```
struct Hideout{ int pos; int slots; };
vector<Hideout> hideouts{{30, 1}, {54, 1}, {123, 1}, {145, 1}, {345, 1}, {423, 1}, {700, 1}};

for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)
{
    auto safety = filter_view(hideouts, [](auto& hideout){ return hideout.slots > 0;});
    auto closest = adjacent_find(safety.begin(), safety.end(), [&](auto& first, auto& second) {
        return (first.pos < fish.pos && second.pos >= fish.pos));

    auto& selected = choose_closer_neighbor(closest);

    swim_for(std::abs(selected->pos - fish.pos) * fish.velocity);

    if(selected_hideout->take_slot())
        stay_hidden();
    else
        try_mimicry();
});
```


Again with the presence of sync it can't be any unseq

We go for par by securing access to slots by mutex

Look for a hideout

```
struct Hideout{ int pos; int slots_cnt(); bool take_slot(); private: int slots; mutex slot_mutex; };
vector<Hideout> hideouts{{30, 1}, {54, 1}, {123, 1}, {145, 1}, {345, 1}, {423, 1}, {700, 1}};

for_each(execution::par, fishes.begin(), fishes.end(), [&](auto& fish)
{
    auto safety = filter_view(hideouts, [](auto& hideout){ return hideout.slots_cnt();
    auto closest = adjacent_find(safety.begin(), safety.end(), [&](auto& first, auto& second) {
        return (first.pos < fish.pos && second.pos >= fish.pos));

    auto& selected = choose_closer_neighbor(closest);

    swim_for(std::abs(selected->pos - fish.pos) * fish.velocity);

    if(selected_hideout->take_slot())
        stay_hidden();
    else
        try_mimicry();
}));
```

one honey bee hive contains 60,000 to 80,000 bees

Become a leader

```
bool has_leader = false;
mutex leaders_mutex;
vector<Fish> fishes;

for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)
{
    unique_lock lock(leaders_mutex);
    while(has_leader){
        lock.unlock(); swim_around(); lock.lock();
    }
    has_leader = true;
    lock.unlock();

    live_leaders_life();

    lock.lock();
    has_leader = false;
});
```

Become a leader

```
bool has_leader = false;
mutex leaders_mutex;
vector<Fish> fishes;
```

```
for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)
{
    unique_lock lock(leaders_mutex);
    while(has_leader){
        lock.unlock(); swim_around(); lock.lock();
    }
    has_leader = true;
    lock.unlock();

    live_leaders_life();

    lock.lock();
    has_leader = false;
});
```

```
unique_lock lock(leaders_mutex);
unique_lock lock(leaders_mutex);
unique_lock lock(leaders_mutex);
unique_lock lock(leaders_mutex);
```

**UNDEFINED
BEHAVIOR!**

Become a leader

```
atomic<bool> has_leader = false;  
vector<Fish> fishes;
```

```
for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)  
{  
    while(has_leader.exchange(true))  
        swim_around();  
  
    live_leaders_life();  
  
    has_leader = false;  
});
```

Become a leader

Nope

First fish will take a leadership and lives it's privileged life.
Second fish will attempt to take the leadership before the first loose it.
Pitty they live in the same thread
Second one will wait forever as the loosing of leadership is sequenced after
Deadlock my friends

```
atomic<bool> has_leader = false;  
vector<Fish> fishes;
```

```
for_each(execution::par_unseq, fishes.begin(), fishes.end(), [&](auto& fish)  
{  
    while(has_leader.exchange(true))  
        swim_around();  
  
    live_leaders_life();  
  
    has_leader = false;  
});
```

```
while(has_leader.exchange(true))  
    swim_around();  
  
live_leaders_life();  
  
while(has_leader.exchange(true))  
    swim_around();
```

DEADLOCK!

Summary

Use standard library algorithms – your life will be easier

Use the strongest possible policy

Try to avoid data dependencies – synchronization impact is huge

Always measure your performance

Beware of non-associative operations

Avoid synchronization points by utilizing views

Deeper insight

Think Parallel (by Bryce Adelstein Lelbach)

The C++ Execution Model (by Bryce Adelstein Lelbach)

Portable floating-point calculations (by Guy Davidson)

Parallel STL <https://github.com/llvm/llvm-project/tree/main/pstl>

Thrust <https://github.com/NVIDIA/cccl/tree/main/thrust>

Sources

All examples and naive algorithms as presented

https://github.com/EmptySquareBubble/Stories_from_a_parallel_universe



Thank you!